

实验二——MapReduce 分布式数据处理

姓名：陈朴炎

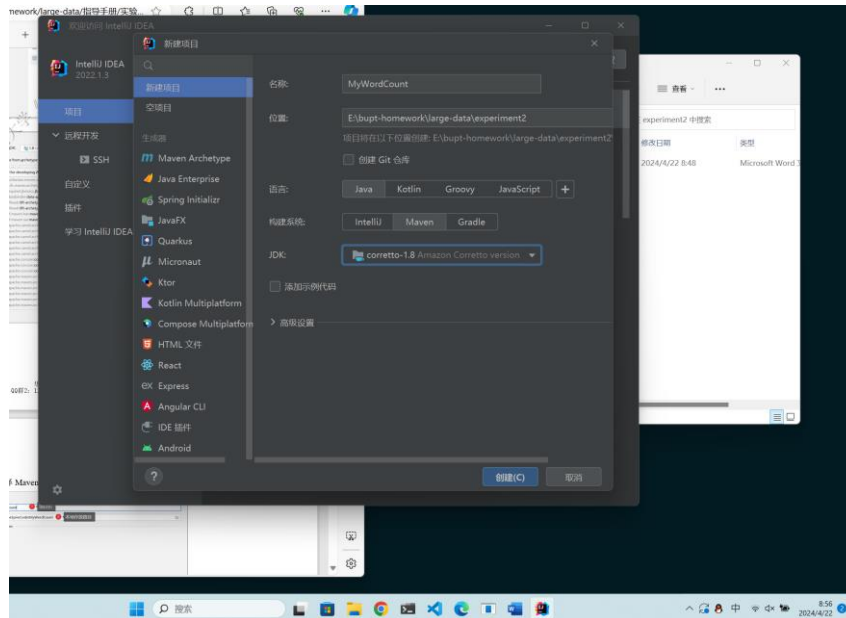
学号：2021211138

目录

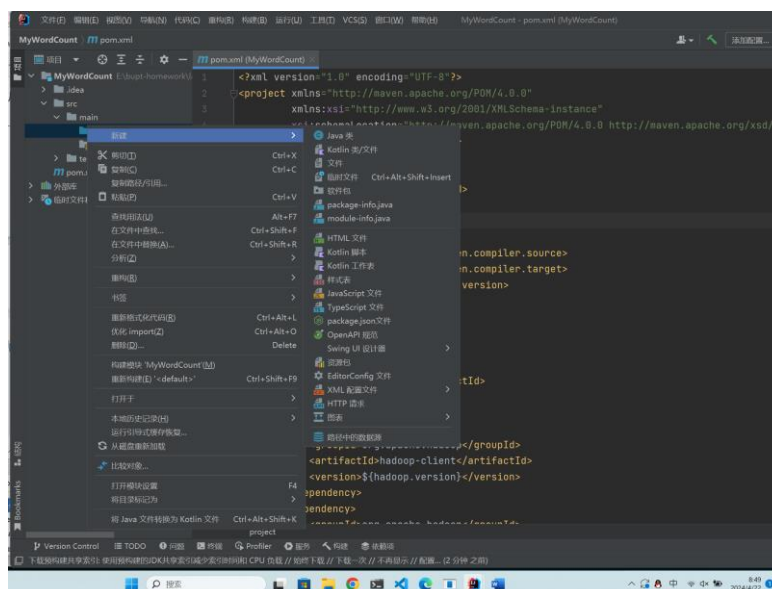
实验二——MapReduce 分布式数据处理	1
1 实验流程	2
2 实验结果与分析	18
2.1 jar 包生成图	18
2.2 执行命令截图	18
2.3 步骤 8 执行命令结果图	19
2.4 步骤 9 执行命令结果图	20
2.5 提交 jar 包（另附件）	21
2.6 解释 WordCount 程序代码	21

1 实验流程


使用 IDEA 新建一个 java 项目，jdk 选择 1.8



新建一个 MyWordCount 项目，修改里面的 pom.xml 文件，将实验一的配置信息复制到这里面。并新建一个 java 类名为 WordCount



修改完 pom.xml 配置文件后，重新构建项目



WordCount 类中添加如下代码

这个类的作用是将输入的文本数据切分成单词，并为每个单词输出一个键值对，其中键是单词本身，值是 1，表示该单词出现了一次。

```
1. public class WordCount {
2.     public static class TokenizerMapper extends Mapper<Object, Text,
        Text, IntWritable>{
3.         private final static IntWritable one = new IntWritable(1);
4.         private Text word = new Text();
5.         public void map(Object key, Text value, Context context)
6.             throws IOException, InterruptedException {
7.             StringTokenizer itr = new StringTokenizer(value.toString(
            ));
8.             while (itr.hasMoreTokens()) {
9.                 word.set(itr.nextToken());
10.                context.write(word, one);
11.            }
12.        }
13.    }
14.
15. }
```

但是很可惜，会有错误，原因是 Text 类在很多包里面都有，项目没办法自己选择所需要的包，需要我们手动添加。

```
    }  
    public class TokenizerMapper extends Mapper<Object, Text, Text, IntWritable> {  
        private static IntWritable one = new IntWritable(1);  
        private Text word = new Text();  
        @Override  
        public void map(Object key, Text value, Context context)  
            throws IOException, InterruptedException {  
            // ...  
        }  
    }  
}
```

导入以下两个包后红色报错消失

```
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Mapper;
```

之后再创建一个 Reducer 对象，继承 hadoop 的 Reducer，用来将 Mapper 的输出结构进行合并和总结，如下图所示。

```
将 Mapper 部分输出的中间结果进行合并和汇总  
public static class IntSumReducer extends Reducer<Text, IntWritable, Text, IntWritable> {  
    private IntWritable result = new IntWritable();  
  
    将 Mapper 部分输出的中间结果进行合并和汇总。  
    形参: key - 当前要计算的键  
    values - 输入值的迭代器  
    context - 上下文对象，用于将结果输出  
    在 reduce 方法内部，首先定义了一个整型变量 sum，用于保存当前键对  
    应的值的总和  
    然后，通过一个 for 循环遍历输入的值的迭代器，将每个值取出并累加到  
    sum 变量中。  
    最后，将 sum 设置到 IntWritable 对象 result 中，并通过上下文对象  
    context 输出该键以及对应的总和。  
  
    public void reduce(Text key, Iterable<IntWritable> values, Context context)  
        throws IOException, InterruptedException {  
        int sum = 0;  
        for (IntWritable val : values) {  
            sum += val.get();  
        }  
        result.set(sum);  
        context.write(key, result);  
    }  
}
```

类和函数说明：

在 reduce 函数作为 Reducer 类的主要函数，它的功能是将 Mapper 部分输出的中间结果进行合并和汇总。在 reduce 函数中，首先定义了一个整型变量 sum，用于保存当前键对应的值的总和。然后，通过一个 for 循环遍历输入的值的迭代器，将每个值取出并累加到 sum 变量中。最后，将 sum 设置到

IntWritable 对象 result 中，并通过上下文对象 context 输出该键以及对应的总和。

上面两个类编写完成之后，可以开始写 WordCount 的主函数了。

主函数的运行逻辑为：

1. 判断输入的参数格式是否正确
2. 创建 job 对象，并为 job 对象初始化
3. 设置输出的 Key 和 Value
4. 设置输入输出的路径

在主函数中，需要先 new 一个 Configuration 对象。需要注意的是，当前这个 Configuration 对象应该选择 apache.hadoop.conf.Configuration 类。

在第四步的时候，要注意 FileInputFormat 和 FileOutputFormat 应该选择导入的包是以下这两个，不然会报错的。

```
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

main 函数如下所示：

```
1.    public static void main (String [] args) throws Exception{
2.
3.        Configuration conf = new Configuration();
4.        // 限定输出参数必须为 2 个
5.        String []otherArgs = new GenericOptionsParser(conf, args).get
        RemainingArgs();
6.
7.        if(otherArgs.length!=2){
8.            System.err.println("Please usage: word-
            count <in> <out>");
9.            System.exit(2);
10.        }
11.
12.        // 创建 job 对象
13.        Job job = new Job(conf, "word count");
```

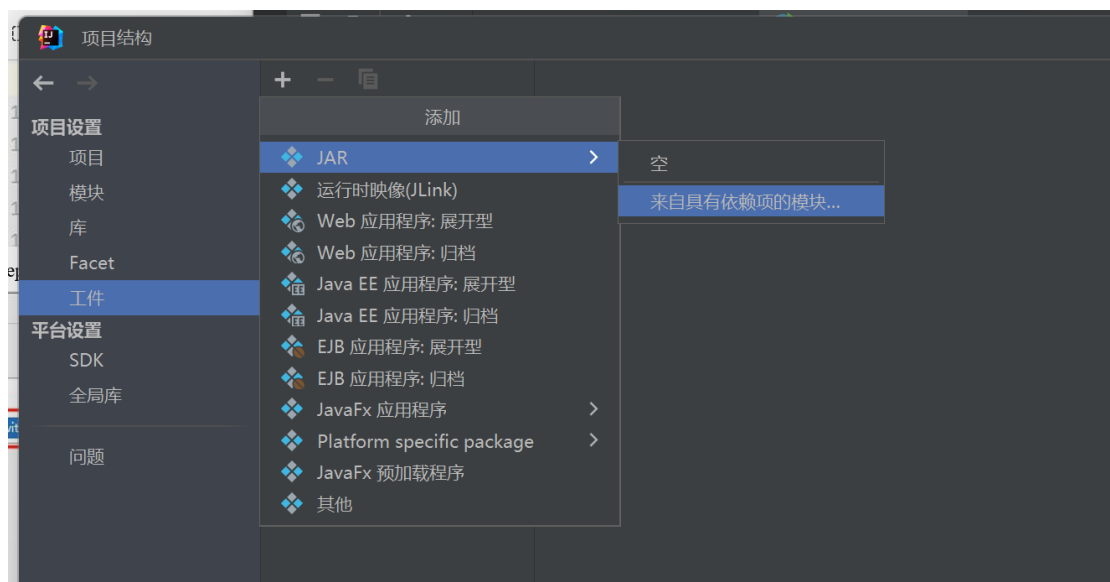
```

14.
15.     // 初始化 job 对象
16.     job.setJarByClass(WordCount.class);
17.     job.setMapperClass(TokenizerMapper.class);
18.     job.setCombinerClass(IntSumReducer.class);
19.     job.setReducerClass(IntSumReducer.class);
20.
21.     // 设置输出格式
22.     job.setOutputKeyClass(Text.class);
23.     job.setOutputValueClass(IntWritable.class);
24.
25.     // 设置输入输出
26.     FileInputFormat.addInputPath(job, new Path(otherArgs[0]));
27.     FileOutputFormat.setOutputPath(job, new Path(otherArgs[1]));
28.
29.     System.exit(job.waitForCompletion(true)?0:1);
30. }

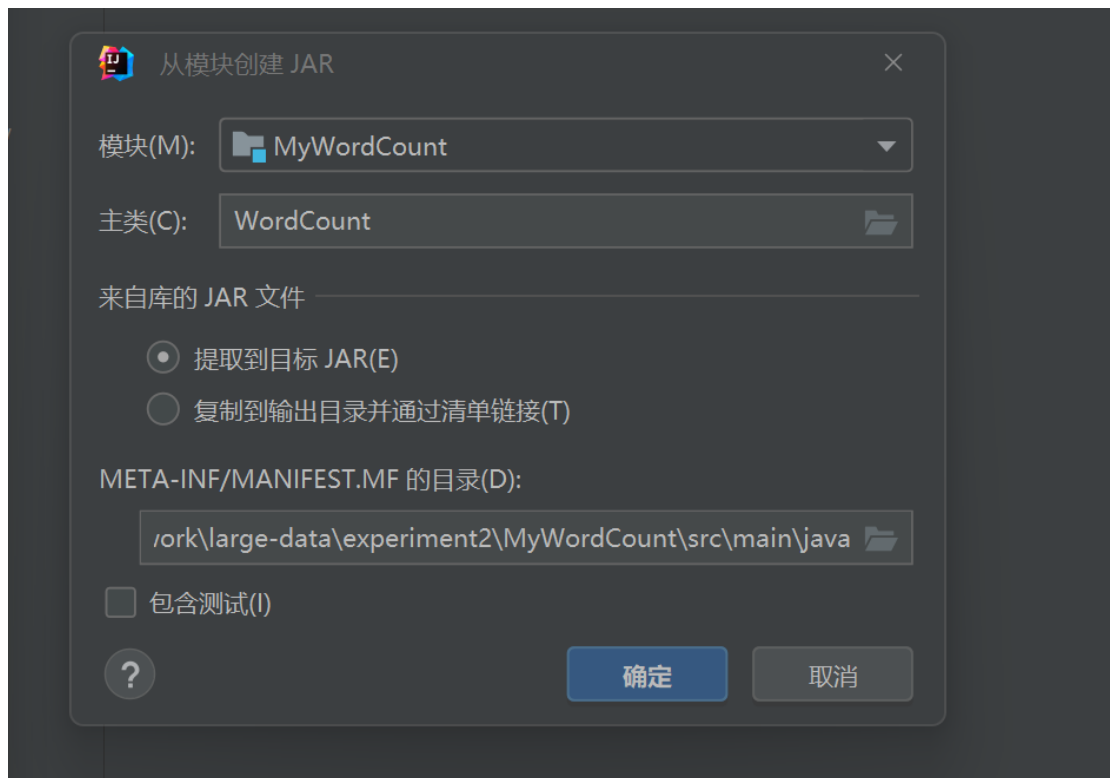
```

接下来应该要打包项目并运行了。

在文件->项目结构中，选择工件，添加，如下图所示

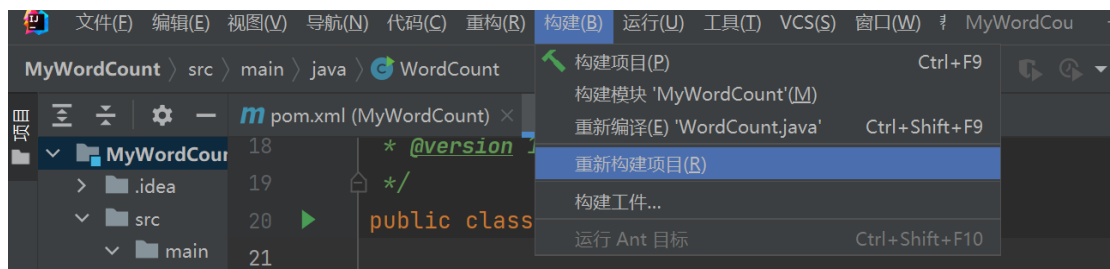


并将主类设置为 WordCount，如下图



之后选择应用，确定即可。


选择 构建 -> 重新构建项目



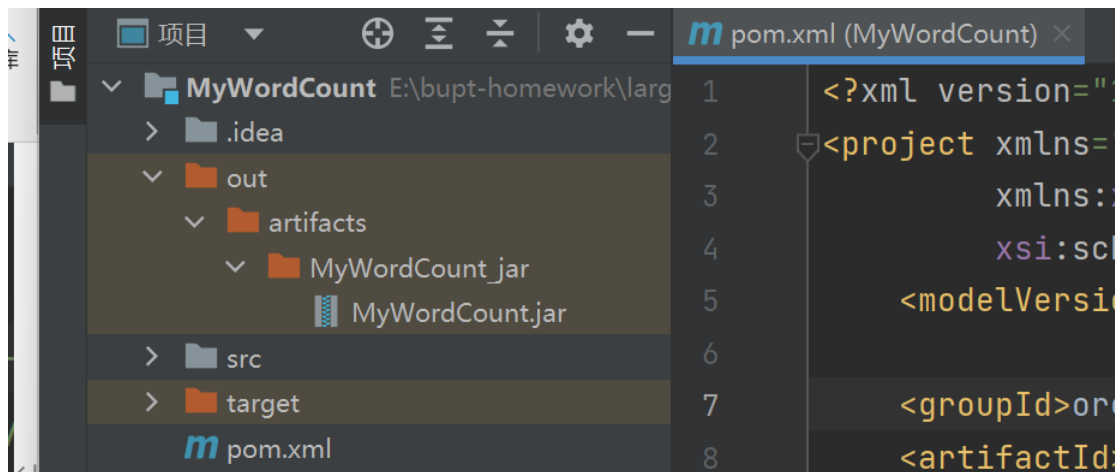
但是这里我没能成功出现 out 包。

原因是没有构建工件。

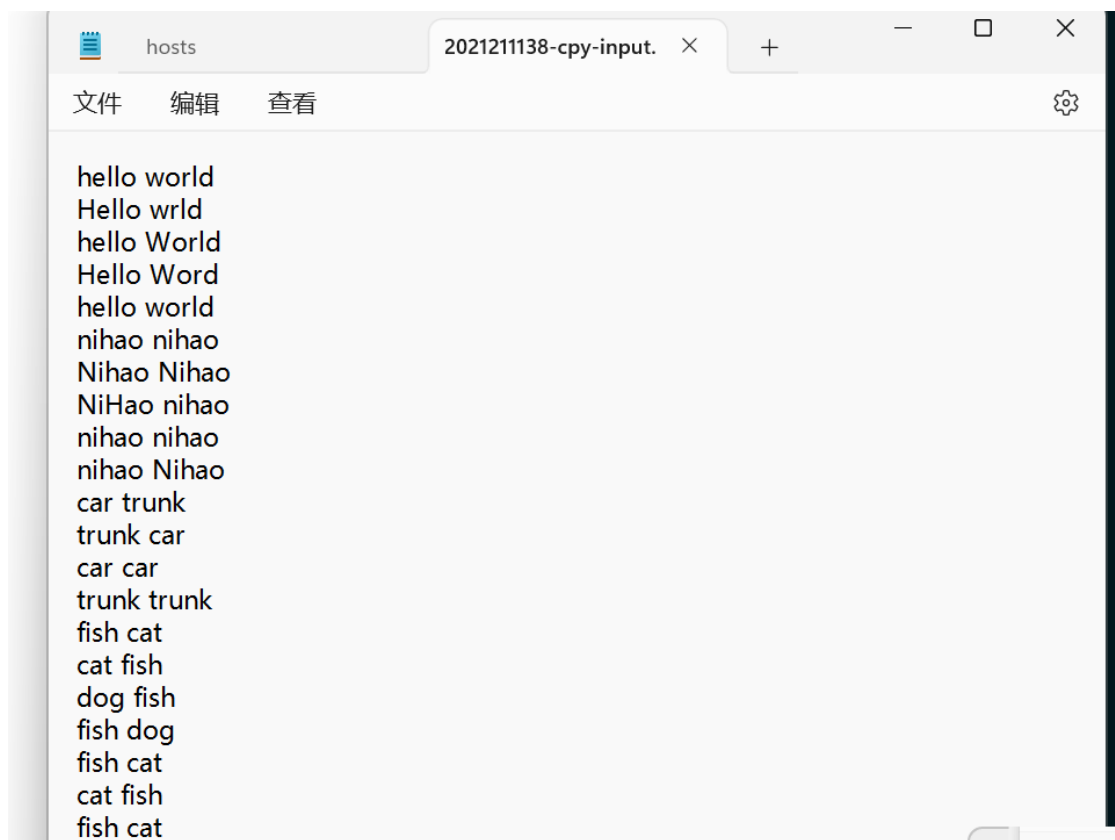

重新构建工件，如下图所示



选择构建工件 -> MyWordCount -> 构建，就出现 out 包了



使用 scp 命令将这个 jar 文件发给服务器



同样的，把这个输入文件传送给服务器


```
E:\bupt-homework\large-data\experiment2\MyWordCount>scp ./2021211138-cpy-input.txt root@1.92.114.12:~/
root@1.92.114.12's password:
2021211138-cpy-input.txt                                100% 210   20.5KB/s   00:00

E:\bupt-homework\large-data\experiment2\MyWordCount>
```

连上四台服务器，启动 hadoop 和 yarn

```
hadoop-2.7.7      javasharedresources  OpenJDK64-BitArchLinux_OpenJ9_0.20.0.tar
[root@cpy-2021211138 ~]# start-dfs.sh start-yarn.sh
Usage: start-dfs.sh [-upgrade|-rollback] [other options such as -clusterId]
[root@cpy-2021211138 ~]# start-dfs.sh
24/04/22 13:24:06 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-j
va classes where applicable
Starting namenodes on [node1]
node1: starting namenode, logging to /home/modules/hadoop-2.7.7/logs/hadoop-root-namenode-cpy-2021211138.out
node2: starting datanode, logging to /home/modules/hadoop-2.7.7/logs/hadoop-root-datanode-cpy-2021211138.out
node3: starting datanode, logging to /home/modules/hadoop-2.7.7/logs/hadoop-root-datanode-cpy-2021211138.out
node4: starting datanode, logging to /home/modules/hadoop-2.7.7/logs/hadoop-root-datanode-cpy-2021211138.out
Starting secondary namenodes [node1]
node1: starting secondarynamenode, logging to /home/modules/hadoop-2.7.7/logs/hadoop-root-secondarynamenode-cpy-2021211
38.out
24/04/22 13:24:21 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-j
va classes where applicable
```

```
[root@cpy-2021211138 ~]# start-yarn.sh
starting yarn daemons
starting resourcemanager, logging to /home/modules/hadoop-2.7.7/logs/yarn-root-resourcemanager-cpy-2021211138.out
node2: starting nodemanager, logging to /home/modules/hadoop-2.7.7/logs/yarn-root-nodemanager-cpy-2021211138.out
node3: starting nodemanager, logging to /home/modules/hadoop-2.7.7/logs/yarn-root-nodemanager-cpy-2021211138.out
node4: starting nodemanager, logging to /home/modules/hadoop-2.7.7/logs/yarn-root-nodemanager-cpy-2021211138.out
[root@cpy-2021211138 ~]#
```

尝试运行 jar 包，主函数抛出异常

```
node4: starting nodemanager, logging to /home/modules/hadoop-2.7.7/logs/yarn-root-nodemanager-cpy-2021211138.
[root@cpy-2021211138 ~]# hadoop jar MyWordCount.jar ./2021211138-cpy-input.txt ./2021211138-cpy-out.txt
Exception in thread "main" java.lang.ClassNotFoundException: ..2021211138-cpy-input.txt
    at java.lang.Class.forNameImpl(Native Method)
    at java.lang.Class.forName(Class.java:402)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:219)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:141)
[root@cpy-2021211138 ~]#
```

我觉得应该是 hdfs 找不到文件导致的

```
mkdir: experiment2: No such file or directory
[root@cpy-2021211138 ~]# hdfs dfs -ls /
24/04/22 13:30:37 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your pla
tform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 3 root supergroup 59 2024-03-22 22:08 /gby_2021211138.txt
-rw-r--r-- 3 root supergroup 37 2024-03-22 22:08 /upload_2021211138.txt
```

使用 `hdfs dfs -ls /` 查看 hdfs 系统根目录 / 文件夹下的内容，发现并没有我们想要的输入文件。

```
24/04/22 13:33:39 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your pla
tform... using builtin-java classes where applicable
Found 3 items
drwxr-xr-x - root supergroup 0 2024-04-22 13:33 /experiment2
-rw-r--r-- 3 root supergroup 59 2024-03-22 22:08 /gby_2021211138.txt
-rw-r--r-- 3 root supergroup 37 2024-03-22 22:08 /upload_2021211138.txt
```

使用 `hdfs dfs -mkdir /experiment2`，创建一个 hdfs 文件夹

使用 `hadoop fs -copyFromLocal <localSrc> <hdfsDst>` 将本地文件拷到 hdfs 文件系统的路径下

```
[root@cpy-2021211138 ~]# hadoop fs -copyFromLocal ~/2021211138-cpy-input.txt /experiment2/
24/04/22 13:37:01 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platfo
rm... using builtin-java classes where applicable
copyFromLocal: '/experiment2/2021211138-cpy-input.txt': File exists
```

又报出如下错误

```
[root@cpy-2021211138 ~]# hadoop jar MyWordCount.jar /experiment2/2021211138-cpy-input.txt /experime
nt2/2021211138-cpy-out.txt
Exception in thread "main" java.lang.ClassNotFoundException: .experiment2.2021211138-cpy-input.txt
    at java.lang.Class.forNameImpl(Native Method)
    at java.lang.Class.forName(Class.java:402)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:219)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:141)
```

原因是 hadoop jar 命令的格式应该是 hadoop jar <JAR 文件路径> <主类> [参数]，而我的命令中缺少了主类的指定。所以应该将命令改为如下所示。

```
[root@cpy-2021211138 ~]# hadoop jar MyWordCount.jar WordCount /experiment2/2021211138-cpy-input
t /experiment2/2021211138-cpy-out.txt
24/04/22 13:39:40 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your pla
rm... using builtin-java classes where applicable
24/04/22 13:39:41 INFO client.RMProxy: Connecting to ResourceManager at node1/192.168.0.30:8032
24/04/22 13:39:43 INFO input.FileInputFormat: Total input paths to process : 1
24/04/22 13:39:43 INFO mapreduce.JobSubmitter: number of splits:1
24/04/22 13:39:43 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1713763501420_0001
24/04/22 13:39:44 INFO impl.YarnClientImpl: Submitted application application_1713763501420_0001
24/04/22 13:39:44 INFO mapreduce.Job: The url to track the job: http://node1:8088/proxy/applica
tion_1713763501420_0001/
24/04/22 13:39:44 INFO mapreduce.Job: Running job: job_1713763501420_0001
```

但是可以从时间戳上看到卡住了从 39 分到 44 分，很久都没有动静。

一直到 45 分，又有动静了

```
24/04/22 13:39:41 INFO client.RMProxy: Connecting to ResourceManager at node1/192.168.0.30:8032
24/04/22 13:39:43 INFO input.FileInputFormat: Total input paths to process : 1
24/04/22 13:39:43 INFO mapreduce.JobSubmitter: number of splits:1
24/04/22 13:39:43 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1713763501420_0001
24/04/22 13:39:44 INFO impl.YarnClientImpl: Submitted application application_1713763501420_0001
24/04/22 13:39:44 INFO mapreduce.Job: The url to track the job: http://node1:8088/proxy/applicatio
n_1713763501420_0001/
24/04/22 13:39:44 INFO mapreduce.Job: Running job: job_1713763501420_0001
24/04/22 13:45:45 INFO mapreduce.Job: Job job_1713763501420_0001 running in uber mode : false
24/04/22 13:45:45 INFO mapreduce.Job: map 0% reduce 0%
```

但是 map 0% reduce 0%，显然不对

剩下的信息如下图所示

```

24/04/22 13:45:45 INFO mapreduce.Job: Job job_1713763501420_0001 running in uber mode : false
24/04/22 13:45:45 INFO mapreduce.Job: map 0% reduce 0%
24/04/22 13:45:45 INFO mapreduce.Job: Job job_1713763501420_0001 failed with state FAILED due to:
Application application_1713763501420_0001 failed 2 times due to Error launching appattempt_171376
3501420_0001_000002. Got exception: java.net.ConnectException: Call From cpy-2021211138/127.0.0.1
to 127.0.0.1:42819 failed on connection exception: java.net.ConnectException: Connection refused;
For more details see: http://wiki.apache.org/hadoop/ConnectionRefused
    at sun.reflect.GeneratedConstructorAccessor37.newInstance(Unknown Source)
    at sun.reflect.DelegatingConstructorAccessorImpl.newInstance(DelegatingConstructorAccessor
Impl.java:45)
    at java.lang.reflect.Constructor.newInstance(Constructor.java:423)
    at org.apache.hadoop.net.NetUtils.wrapWithMessage(NetUtils.java:792)
    at org.apache.hadoop.net.NetUtils.wrapException(NetUtils.java:732)
    at org.apache.hadoop.ipc.Client.call(Client.java:1480)
    at org.apache.hadoop.ipc.Client.call(Client.java:1413)
    at org.apache.hadoop.ipc.ProtobufRpcEngine$Invoker.invoke(ProtobufRpcEngine.java:229)
    at com.sun.proxy.$Proxy83.startContainers(Unknown Source)
    at org.apache.hadoop.yarn.api.impl.pb.client.ContainerManagementProtocolPBClientImpl.start
Containers(ContainerManagementProtocolPBClientImpl.java:96)
    at org.apache.hadoop.yarn.server.resourcemanager.amlauncher.AMLauncher.run(AMLauncher.java:
250)
    at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1149)
    at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:624)
    at java.lang.Thread.run(Thread.java:823)
Caused by: java.net.ConnectException: Connection refused
    at sun.nio.ch.SocketChannelImpl.checkConnect(Native Method)
    at sun.nio.ch.SocketChannelImpl.finishConnect(SocketChannelImpl.java:716)
    at org.apache.hadoop.net.SocketIOWithTimeout.connect(SocketIOWithTimeout.java:206)
    at org.apache.hadoop.net.NetUtils.connect(NetUtils.java:531)
    at org.apache.hadoop.net.NetUtils.connect(NetUtils.java:495)
    at org.apache.hadoop.ipc.Client$Connection.setupConnection(Client.java:615)
    at org.apache.hadoop.ipc.Client$Connection.setupIOstreams(Client.java:713)
    at org.apache.hadoop.ipc.Client$Connection.access$2900(Client.java:376)
    at org.apache.hadoop.ipc.Client.getConnection(Client.java:1529)
    at org.apache.hadoop.ipc.Client.call(Client.java:1452)
    ... 15 more
. Failing the application.
24/04/22 13:45:45 INFO mapreduce.Job: Counters: 0
[root@cpy-2021211138 ~]#

```

从下面这句话中，能得出，连接到 42819 端口失败了。

```

24/04/22 13:45:45 INFO mapreduce.Job: Job job_1713763501420_0001
failed with state FAILED due to: Application
application_1713763501420_0001 failed 2 times due to Error
launching appattempt_1713763501420_0001_000002. Got exception:
java.net.ConnectException: Call From cpy-2021211138/127.0.0.1
to 127.0.0.1:42819 failed on connection exception:
java.net.ConnectException: Connection refused;

```

之后我使用 netstat 命令，并没有看到 42819 端口打开的信息，我觉得应该是端口没打开导致的。

```

[root@cpy-2021211138 ~]# netstat -tuln | grep 42819
[root@cpy-2021211138 ~]#

```

3. 使用 ss 命令

但是重新尝试了很多次每次的端口都不一样，应该不是端口的问题

我查阅了相关资料，有用的如下：

<https://blog.csdn.net/u014646662/article/details/82890443>

对于我这个问题，应该先配置环境变量

首先输入 `vim ~/.bashrc`，在最后面添加几行，如下

```
# .bashrc

# User specific aliases and functions

alias rm='rm -i'
alias cp='cp -i'
alias mv='mv -i'

# Source global definitions
if [ -f /etc/bashrc ]; then
    . /etc/bashrc
fi

export JAVA_HOME=/usr/java/default
export HADOOP_HOME=~/.hadoop-2.7.7
export HADOOP_MAPRED_HOME=$HADOOP_HOME
export HADOOP_COMMON_HOME=$HADOOP_HOME
export HADOOP_HDFS_HOME=$HADOOP_HOME
export HADOOP_MAPRED_HOME=$HADOOP_HOME
export HADOOP_YARN_HOME=$HADOOP_HOME
export HADOOP_INSTALL=$HADOOP_HOME
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_HOME/lib/native
export HADOOP_CONF_DIR=$HADOOP_HOME/etc/hadoop
export HADOOP_LIBEXEC_DIR=$HADOOP_HOME/libexec
export JAVA_LIBRARY_PATH=$HADOOP_HOME/lib/native:$JAVA_LIBRARY_PATH
export PATH=.:$JAVA_HOME/bin:$HADOOP_HOME/bin:$HADOOP_HOME/sbin:$PATH
export CLASSPATH=$JAVA_HOME/lib:$JRE_HOME/lib:$CLASSPATH
```

然后在 `hadoop-2.7.7/etc/hadoop` 中添加一个 `mapred-site.xml` 文件

文件内容如下：

```

<!-- Put site-specific property overrides in this file. -->

<configuration>
<property>
  <name>yarn.app.mapreduce.am.env</name>
  <value>HADOOP_MAPRED_HOME=${HADOOP_HOME}</value>
</property>
<property>
  <name>mapreduce.map.env</name>
  <value>HADOOP_MAPRED_HOME=${HADOOP_HOME}</value>
</property>
<property>
  <name>mapreduce.reduce.env</name>
  <value>HADOOP_MAPRED_HOME=${HADOOP_HOME}</value>
</property>
<property>
  <name>mapreduce.application.classpath</name>
  <value>
    ${HADOOP_HOME}/etc/hadoop,
    ${HADOOP_HOME}/share/hadoop/common/*,
    ${HADOOP_HOME}/share/hadoop/common/lib/*,
    ${HADOOP_HOME}/share/hadoop/hdfs/*,
    ${HADOOP_HOME}/share/hadoop/hdfs/lib/*,
    ${HADOOP_HOME}/share/hadoop/mapreduce/*,
    ${HADOOP_HOME}/share/hadoop/mapreduce/lib/*,
    ${HADOOP_HOME}/share/hadoop/yarn/*,
    ${HADOOP_HOME}/share/hadoop/yarn/lib/*
  </value>
</property>
</configuration>

```

这个文件是为了配置 hadoop 框架的相关参数的

yarn.app.mapreduce.am.env: 这个属性用于配置 MapReduce 应用程序的 ApplicationMaster 环境变量。在这里, 通过设置 HADOOP_MAPRED_HOME 为 \${HADOOP_HOME}, 将 MapReduce 应用程序的环境变量指向 Hadoop 安装目录。

mapreduce.map.env: 这个属性用于配置 Map Task 执行器的环境变量。同样地, 通过设置 HADOOP_MAPRED_HOME 为 \${HADOOP_HOME}, 将 Map Task 执行器的环境变量指向 Hadoop 安装目录。

mapreduce.reduce.env: 这个属性用于配置 Reduce Task 执行器的环境变量, 同样地, 通过设置 HADOOP_MAPRED_HOME 为 \${HADOOP_HOME}, 将 Reduce Task 执行器的环境变量指向 Hadoop 安装目录。

mapreduce.application.classpath: 这个属性用于配置 MapReduce

应用程序的类路径，以便能够加载所需的类和库文件。在这里，指定了各种 Hadoop 相关的目录和库文件的路径，以确保 MapReduce 任务能够正常运行。

保存后，执行 `source ~/.bashrc` 激活新的环境变量

但是这个时候又报错了，如下：

```
hadoop-2.7.7.tar.gz      OpenJDK8U-jdk_aarch64-linux-openj9_8u292b10-openj9-0.26.0.tar
[root@cpy-2021211138 ~]# hadoop jar MyWordCount.jar WordCount /experiment2/2021211138-cpy-input.txt /
experiment2/2021211138-cpy-out.txt
/root/hadoop-2.7.7/bin/hadoop: line 166: /usr/java/default/bin/java: No such file or directory
[root@cpy-2021211138 ~]# echo $JAVA_HOME
```

找不到 `/usr/java/default/` 路径

我通过如下命令找到了 java 的安装路径

```
[root@cpy-2021211138 ~]# ls -l /usr/bin/java
lrwxrwxrwx 1 root root 22 Nov 30 2019 /usr/bin/java -> /etc/alternatives/java
[root@cpy-2021211138 ~]# ls -l /etc/alternatives/java
lrwxrwxrwx 1 root root 74 Nov 30 2019 /etc/alternatives/java -> /usr/lib/jvm/java-1.8.0-openjdk-1.8.
0.232.b09-0.el7_7.aarch64/jre/bin/java
[root@cpy-2021211138 ~]#
```

修改了 `~/.bashrc` 的 `JAVA_HOME` 后，还是不行

```
[root@cpy-2021211138 ~]# vi ~/.bashrc
[root@cpy-2021211138 ~]# hadoop jar MyWordCount.jar WordCount /experiment2/2021211138-cpy-input.txt /
experiment2/2021211138-cpy-out.txt
/root/hadoop-2.7.7/bin/hadoop: line 166: /usr/java/default/bin/java: No such file or directory
[root@cpy-2021211138 ~]#
```

尝试修改 `/etc/profile`

```
umask 022

fi

for i in /etc/profile.d/*.sh /etc/profile.d/sh.local ; do
    if [ -r "$i" ]; then
        if [ "${i##*/}" != "$-" ]; then
            . "$i"
        else
            . "$i" >/dev/null
        fi
    fi
done

unset i
unset -f pathmunge

export JAVA_HOME=/usr/lib/jvm/jdk8u292-b10

export HADOOP_HOME=/home/modules/hadoop-2.7.7
export PATH=$JAVA_HOME/bin:$PATH
export PATH=$HADOOP_HOME/bin:$HADOOP_HOME/sbin:$PATH
export HADOOP_CLASSPATH=/home/modules/hadoop-2.7.7/share/hadoop/tools/lib/*:$HADOOP_CLASSPATH
export HADOOP_MAPRED_HOME=$HADOOP_HOME
export HADOOP_COMMON_HOME=$HADOOP_HOME
export HADOOP_HDFS_HOME=$HADOOP_HOME
export HADOOP_MAPRED_HOME=$HADOOP_HOME
export HADOOP_YARN_HOME=$HADOOP_HOME
export HADOOP_INSTALL=$HADOOP_HOME
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_HOME/lib/native
export HADOOP_CONF_DIR=$HADOOP_HOME
export HADOOP_LIBEXEC_DIR=$HADOOP_HOME/libexec
export JAVA_LIBRARY_PATH=$HADOOP_HOME/lib/native:$JAVA_LIBRARY_PATH
export HADOOP_CONF_DIR=$HADOOP_HOME/etc/hadoop
export PATH=.:$JAVA_HOME/bin:$HADOOP_HOME/bin:$HADOOP_HOME/sbin:$PATH
```

使用 `source /etc/profile` 重新激活

现在进入到 HADOOP_HOME 目录中, 如上图我的是/home/modules/hadoop-2.7.7

进入到 etc/hadoop 中

```
-bash: syntax error near unexpected token `;'
[root@cpy-2021211138 etc]# ls
hadoop
[root@cpy-2021211138 etc]# cd hadoop
[root@cpy-2021211138 hadoop]# ls
capacity-scheduler.xml  httpfs-env.sh          mapred-env.sh
configuration.xml       httpfs-log4j.properties mapred-queues.xml.template
container-executor.cfg  httpfs-signature.secret mapred-site.xml
core-site.xml           httpfs-site.xml        slaves
hadoop-env.cmd          kms-acls.xml           ssl-client.xml.example
hadoop-env.sh           kms-env.sh             ssl-server.xml.example
hadoop-metrics2.properties kms-log4j.properties  yarn-env.cmd
hadoop-metrics.properties kms-site.xml           yarn-env.sh
hadoop-policy.xml       log4j.properties      yarn-site.xml
hdfs-site.xml           mapred-env.cmd
```

修改 mapred-site.xml 文件, 添加如下信息

```
<property>
  <name>yarn.app.mapreduce.am.env</name>
  <value>HADOOP_MAPRED_HOME=${HADOOP_HOME}</value>
</property>
<property>
  <name>mapreduce.map.env</name>
  <value>HADOOP_MAPRED_HOME=${HADOOP_HOME}</value>
</property>
<property>
  <name>mapreduce.reduce.env</name>
  <value>HADOOP_MAPRED_HOME=${HADOOP_HOME}</value>
</property>
<property>
  <name>mapreduce.application.classpath</name>
  <value>
    ${HADOOP_HOME}/etc/hadoop,
    ${HADOOP_HOME}/share/hadoop/common/*,
    ${HADOOP_HOME}/share/hadoop/common/lib/*,
    ${HADOOP_HOME}/share/hadoop/hdfs/*,
    ${HADOOP_HOME}/share/hadoop/hdfs/lib/*,
    ${HADOOP_HOME}/share/hadoop/mapreduce/*,
    ${HADOOP_HOME}/share/hadoop/mapreduce/lib/*,
    ${HADOOP_HOME}/share/hadoop/yarn/*,
    ${HADOOP_HOME}/share/hadoop/yarn/lib/*
  </value>
</property>
</configuration>
```

重新执行 jar 指令, 但还是不行。这时候我认真看了一下报错信息

```
24/04/22 13:45:45 INFO mapreduce.Job: Job job_1713763501420_0001 failed with state FAILED due to:
Application application_1713763501420_0001 failed 2 times due to Error launching appattempt_171376
3501420_0001_000002. Got exception: java.net.ConnectException: Call From cpy-2021211138/127.0.0.1
to 127.0.0.1:42819 failed on connection exception: java.net.ConnectException: Connection refused;
For more details see: http://wiki.apache.org/hadoop/ConnectionRefused
at sun.reflect.GeneratedConstructorAccessor27.newInstance(Unknown Source)
```

它其实是 java 连接错误, 在连接到 127.0.0.1 的时候出错了。为什么它会连到 42819 端口呢? 我们的 hadoop 服务的端口号并不是这个。我重新回顾了

一下这个实验和上一个实验，定位了问题所在：hosts 文件有问题。


打开/etc/hosts 文件

```
# 127.0.0.1    localhost        localhost.localdomain  localhost4        localh
# 127.0.0.1    localhost        localhost
# 127.0.0.1    cpy-2021211138  cpy-2021211138

# 192.168.0.30  node1
# 120.46.149.118 node2
# 1.92.86.3     node3
# 120.46.87.42  node4

192.168.0.30    node1
192.168.0.213   node2
192.168.0.161   node3
192.168.0.135   node4

192.168.0.30    node1    node1
# 127.0.0.1     cpy-2021211138  cpy-2021211138
```



因为服务器每次启动都会自己添加一个 127.0.0.1 的配置，所以它这里自动多了一行 127.0.0.1 的配置信息。这也是为什么我始终运行不成功的原因了。

把这一行删掉，或者注释掉之后，重新执行 jar 指令。执行过程及结果如下图所示：

```
[root@cpy-2021211138 ~]# hadoop jar MyWordCount.jar WordCount /experiment2/2021211138-cpy-input.txt /exp
eriment2/2021211138-cpy-output.txt
24/04/22 21:12:45 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform...
using builtin-java classes where applicable
24/04/22 21:12:46 INFO client.RMPProxy: Connecting to ResourceManager at node1/192.168.0.30:8032
24/04/22 21:12:48 INFO input.FileInputFormat: Total input paths to process : 1
24/04/22 21:12:48 INFO mapreduce.JobSubmitter: number of splits:1
24/04/22 21:12:48 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1713791462926_0001
24/04/22 21:12:48 INFO impl.YarnClientImpl: Submitted application application_1713791462926_0001
24/04/22 21:12:48 INFO mapreduce.Job: The url to track the job: http://node1:8088/proxy/application_1713
791462926_0001/
24/04/22 21:12:48 INFO mapreduce.Job: Running job: job_1713791462926_0001
24/04/22 21:13:00 INFO mapreduce.Job: Job job_1713791462926_0001 running in uber mode : false
24/04/22 21:13:00 INFO mapreduce.Job: map 0% reduce 0%
24/04/22 21:13:04 INFO mapreduce.Job: map 100% reduce 0%
24/04/22 21:13:09 INFO mapreduce.Job: map 100% reduce 100%
24/04/22 21:13:10 INFO mapreduce.Job: Job job_1713791462926_0001 completed successfully
24/04/22 21:13:10 INFO mapreduce.Job: Counters: 49
```



```

24/04/22 21:13:10 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=165
    FILE: Number of bytes written=253147
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=359
    HDFS: Number of bytes written=103
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=2409
    Total time spent by all reduces in occupied slots (ms)=2652
    Total time spent by all map tasks (ms)=2409
    Total time spent by all reduce tasks (ms)=2652
    Total vcore-milliseconds taken by all map tasks=2409
    Total vcore-milliseconds taken by all reduce tasks=2652
    Total megabyte-milliseconds taken by all map tasks=2466816
    Total megabyte-milliseconds taken by all reduce tasks=2715648
  Map-Reduce Framework
    Map input records=21
    Map output records=42
    Map output bytes=389
    Map output materialized bytes=165
    Input split bytes=119
    Combine input records=42
    Combine output records=14
    Reduce input groups=14
    Reduce shuffle bytes=165
    Reduce input records=14
    Reduce output records=14
    Spilled Records=28
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=104
    CPU time spent (ms)=820
    Physical memory (bytes) snapshot=338493440
    Virtual memory (bytes) snapshot=2584477696
    Total committed heap usage (bytes)=172490752
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Input Format Counters
    Bytes Read=240
  File Output Format Counters
    Bytes Written=103

```

最后执行命令：

hadoop fs -cat /experiment2/2021211138-cpy-output.txt/part-r-00000, 结果如下：

```
[root@cpy-2021211138 ~]# hadoop fs -cat /experiment2/2021211138-cpy-output.txt/part-r-00000
24/04/22 21:20:53 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform...
using builtin-java classes where applicable
Hello      2
NiHao      1
Nihao      3
Word       1
World      1
car        4
cat        5
dog        2
fish       7
hello      3
nihao      6
trunk      4
world      2
wrld       1
[root@cpy-2021211138 ~]# cd $HADOOP_HOME
```

实验就顺利完成了。

2 实验结果与分析


其实实验结果与分析已经在 [1 实验流程](#)里都有了。这里我再次放一下。

2.1 jar 包生成图



2.2 执行命令截图

1. 发送 jar 包



2. 发送 input 文件

```
E:\bupt-homework\large-data\experiment2\MyWordCount>scp ./2021211138-cpy-input.txt root@192.114.12:~/
root@192.114.12's password:
2021211138-cpy-input.txt                                100% 210   20.5KB/s   00:00

E:\bupt-homework\large-data\experiment2\MyWordCount>
```

3. 启动 hadoop 集群

```
hadoop-2.7.7 javasharedresources OpenJDK64-BitArch Linux OpenJ9_6u292b10_openj9-0.20.0.tar
[root@cpy-2021211138 ~]# start-dfs.sh start-yarn.sh
Usage: start-dfs.sh [-upgrade|-rollback] [other options such as -clusterId]
[root@cpy-2021211138 ~]# start-dfs.sh
24/04/22 13:24:06 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-j
va classes where applicable
Starting namenodes on [node1]
node1: starting namenode, logging to /home/modules/hadoop-2.7.7/logs/hadoop-root-namenode-cpy-2021211138.out
node2: starting datanode, logging to /home/modules/hadoop-2.7.7/logs/hadoop-root-datanode-cpy-2021211138.out
node3: starting datanode, logging to /home/modules/hadoop-2.7.7/logs/hadoop-root-datanode-cpy-2021211138.out
node4: starting datanode, logging to /home/modules/hadoop-2.7.7/logs/hadoop-root-datanode-cpy-2021211138.out
Starting secondary namenodes [node1]
node1: starting secondarynamenode, logging to /home/modules/hadoop-2.7.7/logs/hadoop-root-secondarynamenode-cpy-2021211
38.out
24/04/22 13:24:21 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-j
va classes where applicable
```

4. 查看 hdfs 根目录, 新建目录 experiment2

```
[root@cpy-2021211138 ~]# hdfs dfs -mkdir /experiment2
24/04/22 13:33:39 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your pla
tform... using builtin-java classes where applicable
[root@cpy-2021211138 ~]# hdfs dfs -ls /
24/04/22 13:33:54 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your pla
tform... using builtin-java classes where applicable
Found 3 items
drwxr-xr-x   - root supergroup            0 2024-04-22 13:33 /experiment2
-rw-r--r--   3 root supergroup            59 2024-03-22 22:08 /gby_2021211138.txt
-rw-r--r--   3 root supergroup            37 2024-03-22 22:08 /upload_2021211138.txt
```

5. 将输入文件放入到新建的目录下

```
[root@cpy-2021211138 ~]# hadoop fs -copyFromLocal ~/2021211138-cpy-input.txt /experiment2/
24/04/22 13:37:01 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platfo
rm... using builtin-java classes where applicable
copyFromLocal: '/experiment2/2021211138-cpy-input.txt': File exists
```

2.3 步骤 8 执行命令结果图

第 8 步这里实验指导手册里命令格式漏了一个好像, 应该如下所示:

hadoop jar <jar 包> <main 函数的入口类> <input 文件或路
径> <output 路径>。

```
[root@cpy-2021211138 ~]# hadoop jar MyWordCount.jar WordCount /experiment2/2021211138-cpy-input.txt /exp
eriment2/2021211138-cpy-output.txt
24/04/22 21:12:45 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform...
using builtin-java classes where applicable
24/04/22 21:12:46 INFO client.RMPProxy: Connecting to ResourceManager at node1/192.168.0.30:8032
24/04/22 21:12:48 INFO input.FileInputFormat: Total input paths to process : 1
24/04/22 21:12:48 INFO mapreduce.JobSubmitter: number of splits:1
24/04/22 21:12:48 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1713791462926_0001
24/04/22 21:12:48 INFO impl.YarnClientImpl: Submitted application application_1713791462926_0001
24/04/22 21:12:48 INFO mapreduce.Job: The url to track the job: http://node1:8088/proxy/application_1713
791462926_0001/
24/04/22 21:12:48 INFO mapreduce.Job: Running job: job_1713791462926_0001
24/04/22 21:13:00 INFO mapreduce.Job: Job job_1713791462926_0001 running in uber mode : false
24/04/22 21:13:00 INFO mapreduce.Job: map 0% reduce 0%
24/04/22 21:13:04 INFO mapreduce.Job: map 100% reduce 0%
24/04/22 21:13:09 INFO mapreduce.Job: map 100% reduce 100%
24/04/22 21:13:10 INFO mapreduce.Job: Job job_1713791462926_0001 completed successfully
24/04/22 21:13:10 INFO mapreduce.Job: Counters: 49
```

```

24/04/22 21:13:10 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=165
    FILE: Number of bytes written=253147
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=359
    HDFS: Number of bytes written=103
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=2409
    Total time spent by all reduces in occupied slots (ms)=2652
    Total time spent by all map tasks (ms)=2409
    Total time spent by all reduce tasks (ms)=2652
    Total vcore-milliseconds taken by all map tasks=2409
    Total vcore-milliseconds taken by all reduce tasks=2652
    Total megabyte-milliseconds taken by all map tasks=2466816
    Total megabyte-milliseconds taken by all reduce tasks=2715648

```

```

  Map-Reduce Framework
    Map input records=21
    Map output records=42
    Map output bytes=389
    Map output materialized bytes=165
    Input split bytes=119
    Combine input records=42
    Combine output records=14
    Reduce input groups=14
    Reduce shuffle bytes=165
    Reduce input records=14
    Reduce output records=14
    Spilled Records=28
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=104
    CPU time spent (ms)=820
    Physical memory (bytes) snapshot=338493440
    Virtual memory (bytes) snapshot=2584477696
    Total committed heap usage (bytes)=172490752

```

```

    Total committed heap usage (bytes)=172490752
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Input Format Counters
    Bytes Read=240
  File Output Format Counters
    Bytes Written=103

```

2.4 步骤 9 执行命令结果图

```

[root@cpy-2021211138 ~]# hadoop fs -cat /experiment2/2021211138-cpy-output.txt/part-r-00000
24/04/22 21:20:53 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform...
using builtin-java classes where applicable
Hello      2
NiHao     1
NiHao     3
Word       1
World      1
car        4
cat        5
dog        2
fish       7
hello      3
nihao      6
trunk      4
world      2
wrlld     1
[root@cpy-2021211138 ~]# cd $HADOOP_HOME

```

2.5 提交 jar 包 (另附件)

2.6 解释 WordCount 程序代码

类 TokenizerMapper:这个类的作用是将输入的文本数据切分成单词,并为每个单词输出一个键值对 其中键是单词本身, 值是 1, 表示该单词出现了一次。

```
这个类的作用是将输入的文本数据切分成单词, 并为每个单词输出一个键值对 其中键是单词本身, 值是1, 表示该单词出现了一次。

public static class TokenizerMapper extends Mapper<Object, Text, Text, IntWritable> {
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();
    public void map(Object key, Text value, Context context)
        throws IOException, InterruptedException {
        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one);
        }
    }
}
```

类 IntSumReducer: 将 Mapper 输出的中间结果进行合并和汇总

Reduce 函数:

形参:

key – 当前要计算的键

values – 输入值的迭代器

context – 上下文对象, 用于将结果输出

函数概述:

在 reduce 方法内部, 首先定义了一个整型变量 sum, 用于保存当前键对应的值的总和

然后, 通过一个 for 循环遍历输入的值的迭代器, 将每个值取出并累加到

sum 变量中。

最后，将 sum 设置到 IntWritable 对象 result 中，并通过上下文对象 context 输出该键以及对应的总和。

```
public static class IntSumReducer extends Reducer<Text, IntWritable, Text, IntWritable> {  
    private IntWritable result = new IntWritable();  
  
    将 Mapper 部分输出的中间结果进行合并和汇总。  
    形参: key - 当前要计算的键  
         values - 输入值的迭代器  
         context - 上下文对象，用于将结果输出  
    在 reduce 方法内部，首先定义了一个整型变量 sum，用于保存当前键对应的值的总和  
    然后，通过一个 for 循环遍历输入的值的迭代器，将每个值取出并累加到 sum 变量中。  
    最后，将 sum 设置到 IntWritable 对象 result 中，并通过上下文对象 context 输出该键以及对应的总和。  
  
    public void reduce(Text key, Iterable<IntWritable> values, Context context)  
        throws IOException, InterruptedException {  
        int sum = 0;  
        for (IntWritable val : values) {  
            sum += val.get();  
        }  
        result.set(sum);  
        context.write(key, result);  
    }  
}
```

主函数的运行逻辑为：

1. 判断输入的参数格式是否正确
2. 创建 job 对象，并为 job 对象初始化
3. 设置输出的 Key 和 Value
4. 设置输入输出的路径

```
public static void main (String [] args) throws Exception{

    Configuration conf = new Configuration();
    // 限定输出参数必须为2个
    String []otherArgs = new GenericOptionsParser(conf, args).getRemainingArgs();

    if(otherArgs.length!=2){
        System.err.println("Please usage: word-count <in> <out>");
        System.exit(2);
    }

    // 创建job对象
    Job job = new Job(conf, "word count");

    // 初始化job对象
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);

    // 设置输出格式
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);

    // 设置输入输出
    FileInputFormat.addInputPath(job,new Path(otherArgs[0]));
    FileOutputFormat.setOutputPath(job, new Path(otherArgs[1]));
    System.exit(job.waitForCompletion(true)?0:1);
}
```